

Sync Clickhouse with MySQL/MongoDB

About

Company:
Xiaoxin Tech.

Industry:
Education

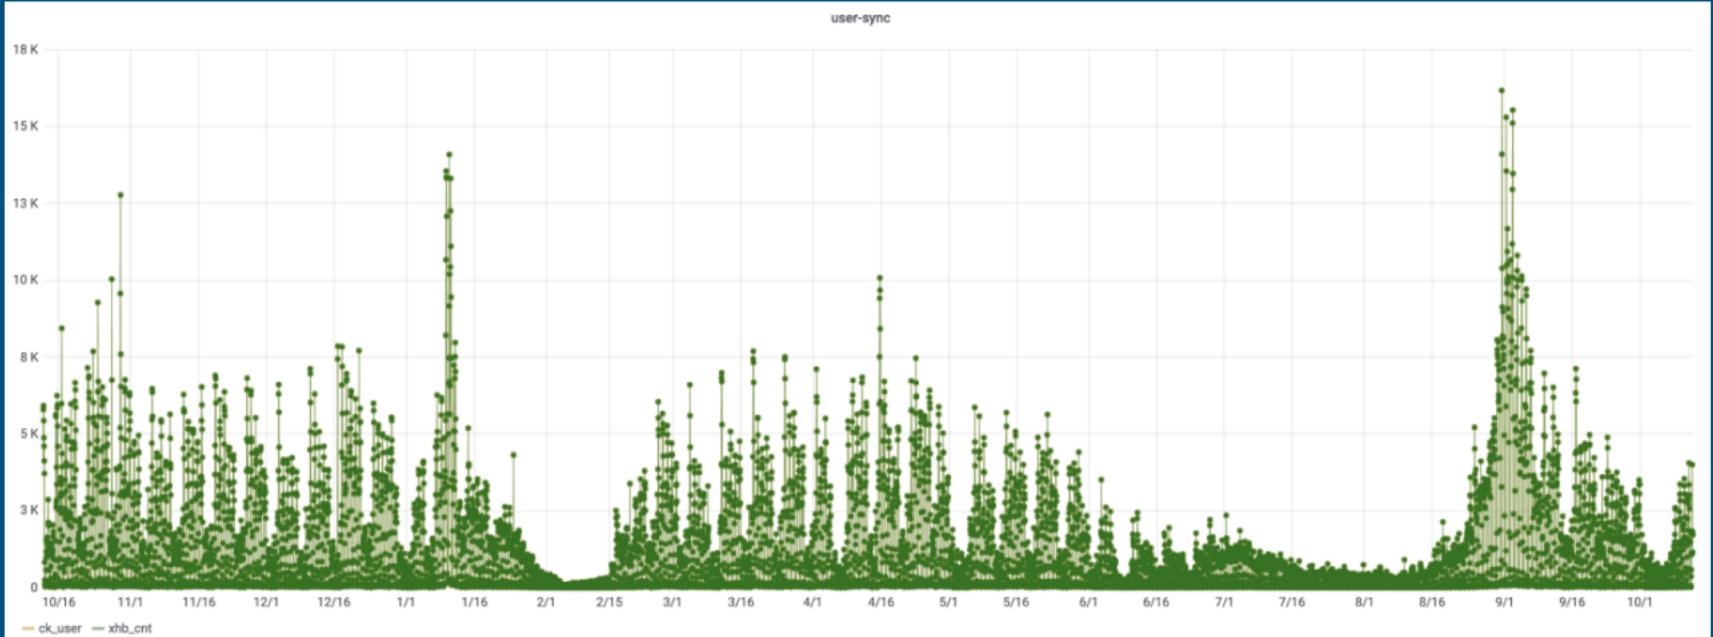
Team:
Big Data

Leader:
wangchao@xiaoheiban.cn

100 billion data this year till now

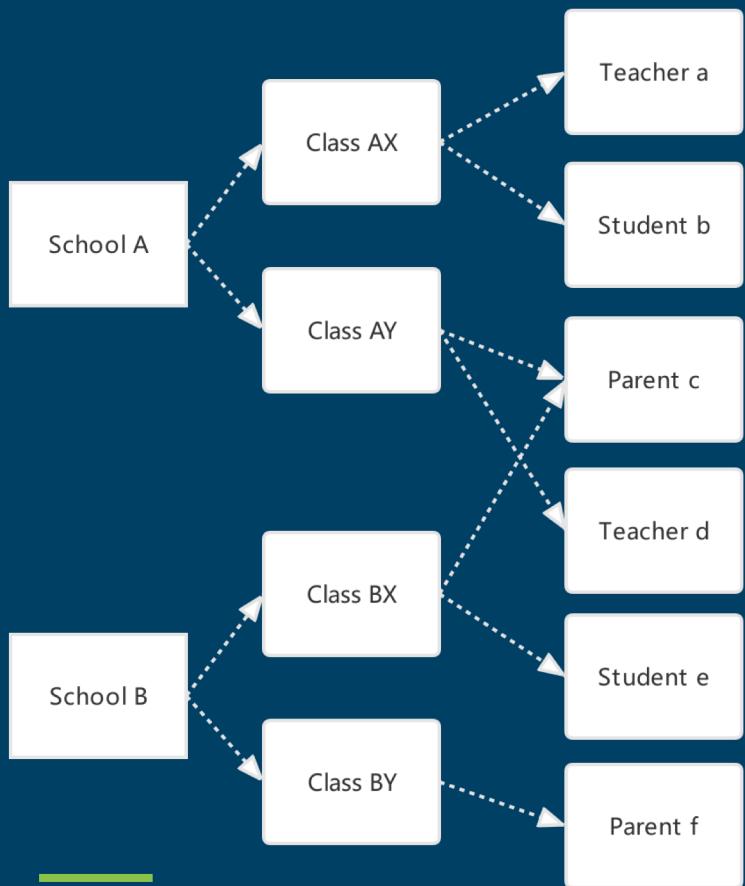


30 million users



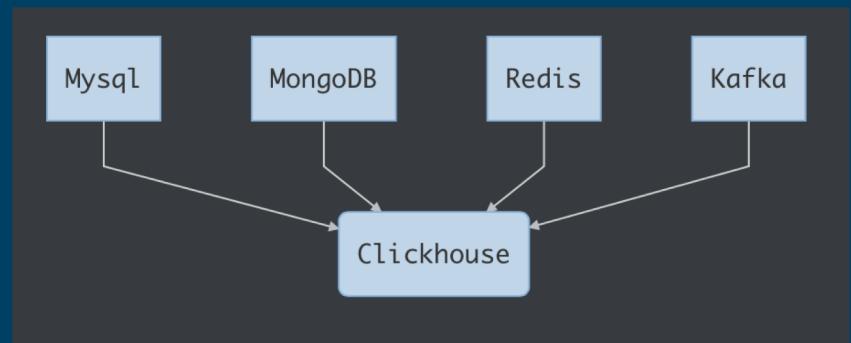
We use Clickhouse in our
daily tasks

Challenges



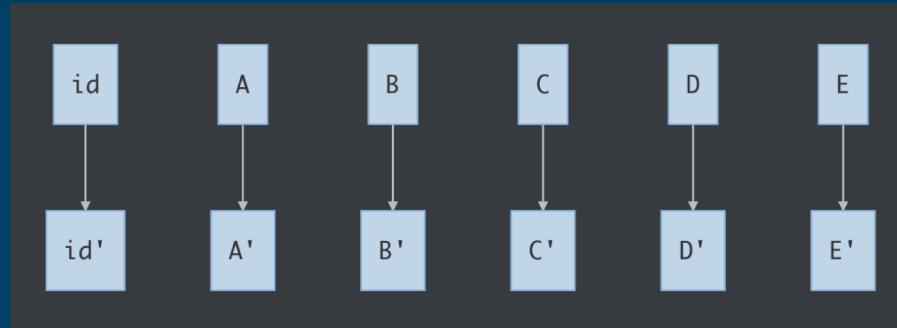
Channllenges

Complex Datasource



Channllenges

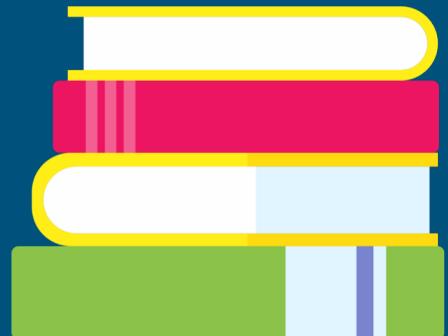
Frequent Updates



Possible Solutions 1.

Replay binlog/oplog CRUD directly

Can't update/delete table frequently in Clickhouse



Possible Solutions 2.

MySQL Engine

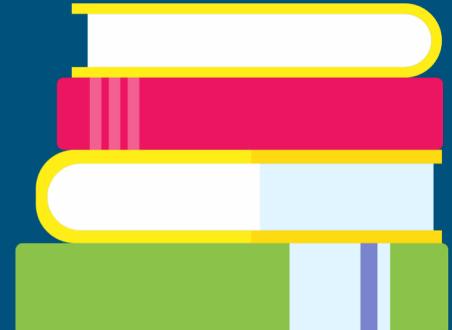
Not suitable for big tables

Not suitable for MongoDB

≡ MySQL

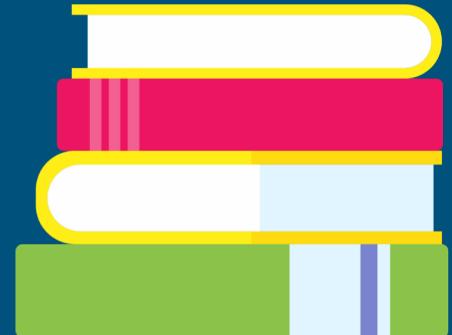
Creating a Database ¶

```
CREATE DATABASE [ IF NOT EXISTS] db_name [ON CLUSTER cluster]
ENGINE = MySQL('host:port', 'database', 'user', 'password')
```



Possible Solutions 3.

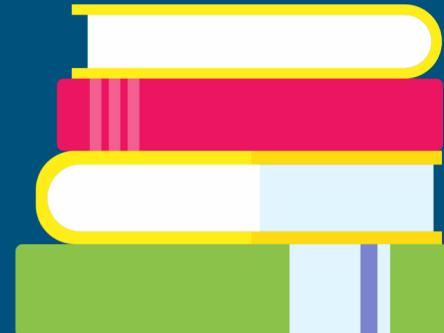
Reinit whole table every day.....



Possible Solutions 4.

CollapsingMergeTree

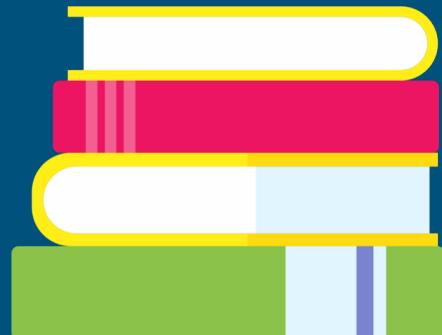
- FINAL is slow
- GROUP BY id HAVING sum(sign)>0
 - Need to use GROUP BY in every query
 - Not suitable for multi-column primary key



Our Solution: PTS

Key Features

- Only one config file needed for a new Clickhouse table
- Init and keep syncing data in one app for a table
- Sync multiple data source to Clickhouse in minutes



PTS

Provider Transform Sinker

- Major Provider Must Listen
- Merge sharding tables
- SolidKey

Typical Provider Config

{

```
Type: mysql,          // mysql, mongodb, redis  
Listen: binlog, // binlog, kafka  
DataSource: user:pass@tcp(example.com:3306)/user,
```

```
Table: user,
```

```
QueryKeys: [ // usually primary key
```

```
    id
```

```
],
```

```
Pairs: { // field mapping
```

```
    id: id,
```

```
    name: name
```

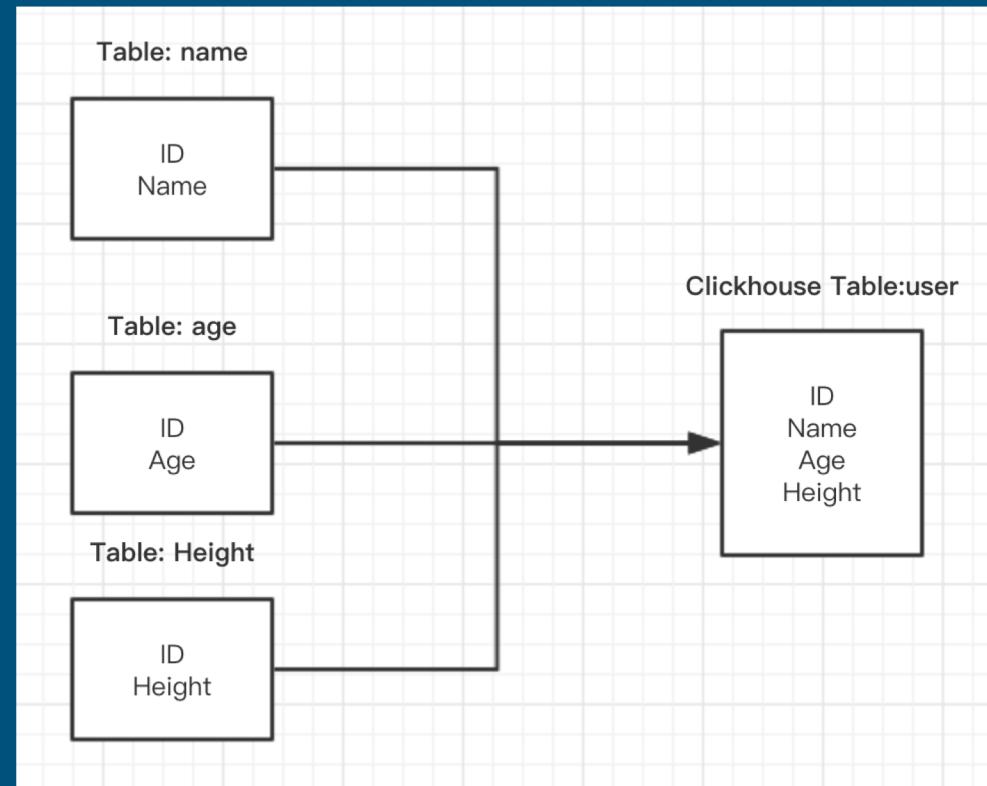
```
}
```

```
}
```

PTS

Provider Transform Sinker

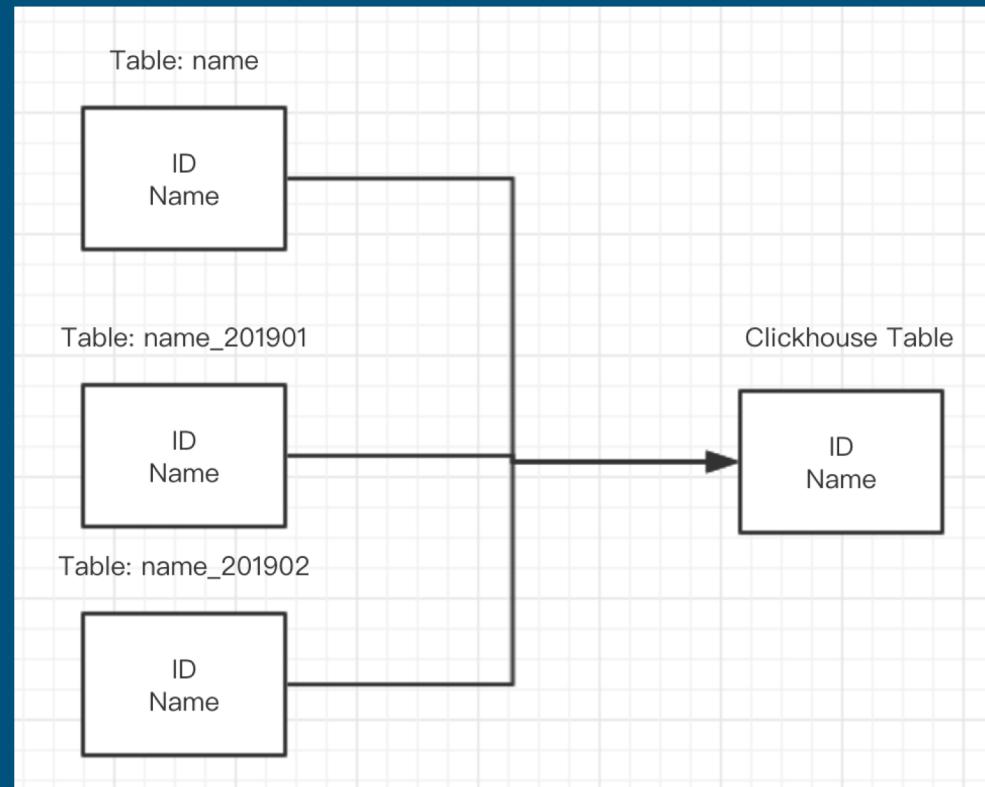
- Major Provider
- Secondary Providers



PTS

Provider Transform Sinker

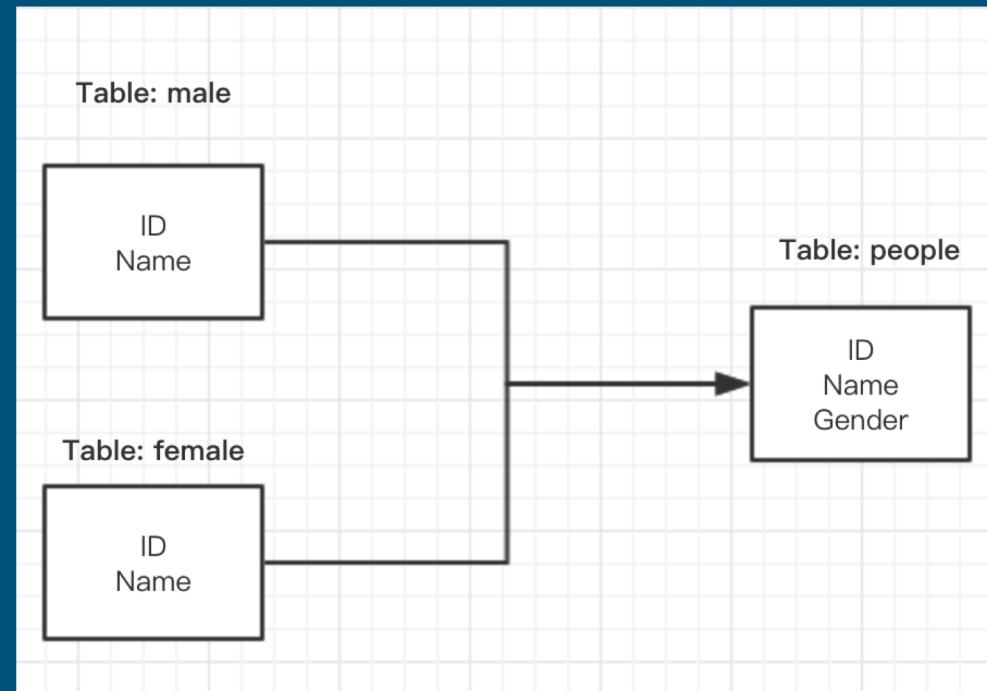
- Merge sharding tables



PTS

Provider Transform Sinker

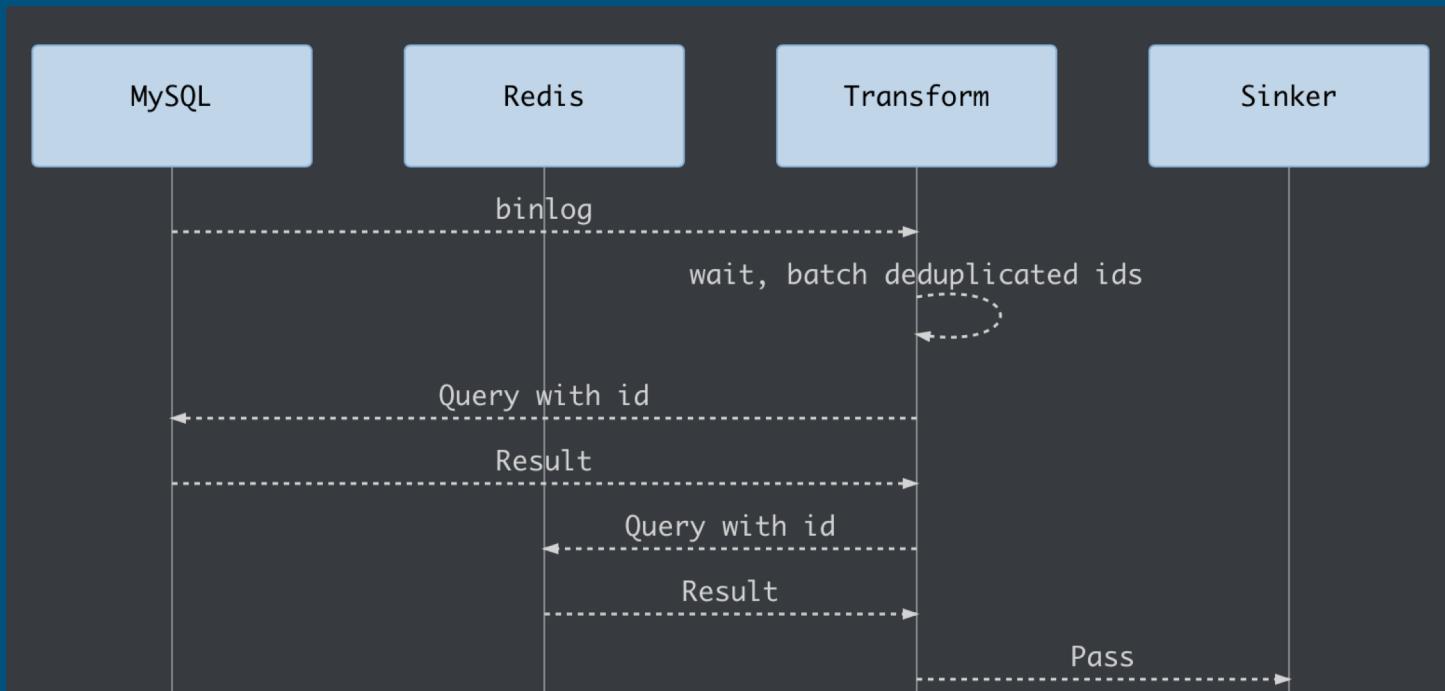
- SolidKey



PTS

Provider Transform Sinker

Transform



PTS

Provider Transform Sinker

- Batch Inserts
 - Time trigger
 - Number trigger
- Magical Flag

PTS

Magical Flag

Quiz



```
INSERT INTO classroom_all(id, ... ,flag)
VALUES (2, ... ,0);
```

```
ALTER TABLE classroom_all
UPDATE flag=flag+1
WHERE (flag=0 or flag=1)
AND id IN (2, ... );
```

PTS

Magical Flag

Create



```
INSERT INTO classroom_all(id, ... ,flag)
VALUES (2, ... ,0);
```

```
ALTER TABLE classroom_all
UPDATE flag=flag+1
WHERE (flag=0 or flag=1)
AND id IN (2, ... );
```

PTS

Magical Flag

Create

id	name	flag
2	Bob1	0



id	name	flag
2	Bob1	1

PTS

Magical Flag

Read



```
SELECT * FROM classroom_all WHERE flag=1
```

PTS

Magical Flag

Update

```
● ● ●  
INSERT INTO classroom_all(id, ... ,flag)  
VALUES (2, ... ,0);
```

```
ALTER TABLE classroom_all  
UPDATE flag=flag+1  
WHERE (flag=0 or flag=1)  
AND id IN (2, ... );
```

PTS

Magical Flag

—
Update

id	name	flag
2	Bob1	1
2	Bob2	0



id	name	flag
2	Bob1	2
2	Bob2	1

PTS

Magical Flag

Delete



```
INSERT INTO classroom_all(id, ... ,flag) VALUES (2 ... ,3);

ALTER TABLE classroom_all
UPDATE flag=flag+1
WHERE flag=1
AND id IN (2 ... );
```

PTS

Magical Flag

Delete

id	name	flag
2	Bob1	2
2	Bob2	1
2	Bob2	3



id	name	flag
2	Bob1	2
2	Bob2	2
2	Bob2	3

PTS

Magical Flag

Add new column

```
SELECT id FROM source_table;  
INSERT INTO classroom_all(id, ... ,flag) VALUES (2 ... ,5);
```

```
ALTER TABLE classroom_all  
UPDATE flag=(flag+3)%7  
WHERE flag=1  
AND id IN (2 ... );
```

```
ALTER TABLE classroom_all  
DELETE WHERE id IN (2 ... );
```

PTS

Magical Flag

Why fast?

Clickhouse Mutation

```
MergeTreeData ::MutableDataPartPtr MergeTreeDataMergerMutator ::mutatePartToTemporaryPart
...
NameSet files_to_skip = {"checksums.txt", "columns.txt"}; //files to skip hardlink
...
for (const auto & entry : in_header) // mutation fields, only flag
{
    IDataType :: StreamCallback callback = [&](const IDataType :: SubstreamPath & substream_path)
    {
        String stream_name = IDataType :: getFileNameForStream(entry.name, substream_path);
        files_to_skip.insert(stream_name + ".bin");
        files_to_skip.insert(stream_name + mrk_extension);
    };
    IDataType :: SubstreamPath stream_path;
    entry.type->enumerateStreams(callback, stream_path);
}
for (Poco :: DirectoryIterator dir_it(source_part->getFullPath()); dir_it != dir_end; ++dir_it)
{
    if (files_to_skip.count(dir_it.name()))
        continue;
    Poco :: Path destination(new_part_tmp_path);
    destination.append(dir_it.name());
    createHardLink(dir_it.path().toString(), destination.toString());
}
```

PTS

Magical Flag

Why fast?

Clickhouse Mutation

```
MergedColumnOnlyOutputStream out(
    data,
    in_header,
    new_part_tmp_path,
    /* sync = */ false,
    compression_codec,
    /* skip_offsets = */ false,
    unused_written_offsets,
    source_part->index_granularity,
    &source_part->index_granularity_info
);

in->readPrefix();
out.writePrefix();

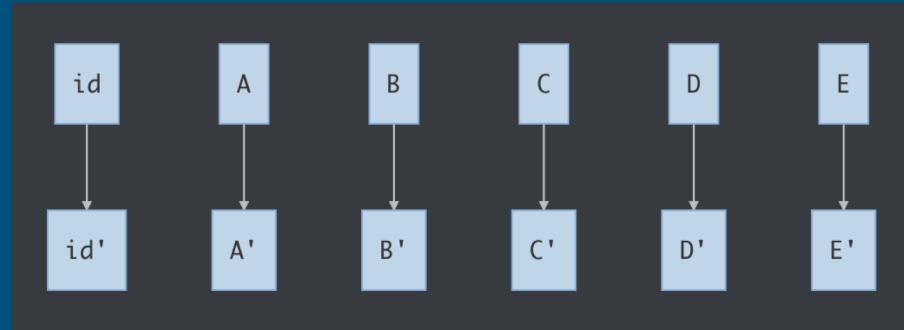
Block block;
while (check_not_cancelled() && (block = in->read()))
{
    out.write(block);

    merge_entry->rows_written += block.rows();
    merge_entry->bytes_written_uncompressed += block.bytes();
}
```

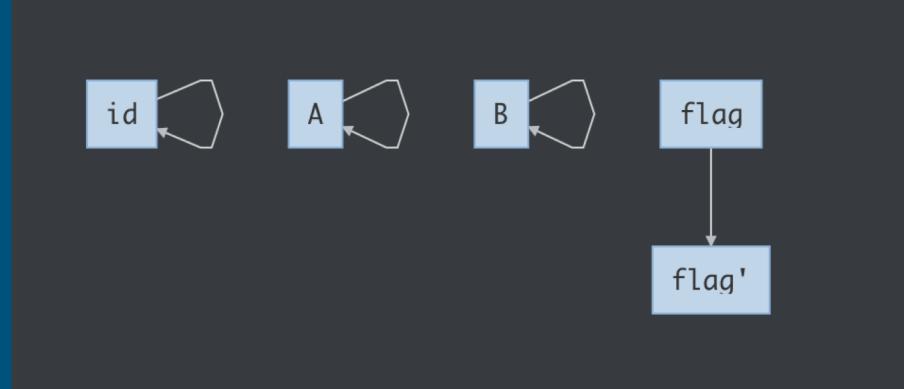
PTS

Magical Flag

Why fast?

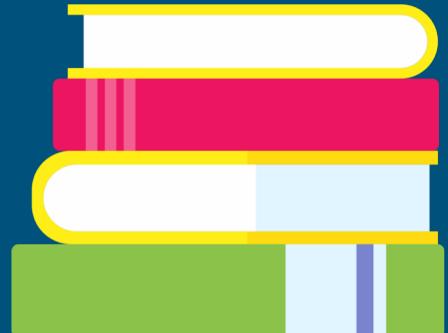


Clickhouse Mutation



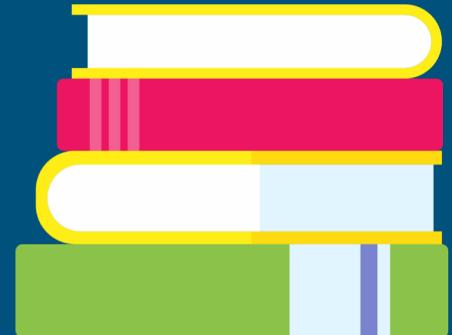
Trouble Shoot

- SQL is too big(max_ast_elements, max_expanded_ast_elements)
- Mutations are stuck (KILL MUTATION)
- Zookeeper OOM because of SQL length (Put ids in a Memory Engine temp table)



Final Product

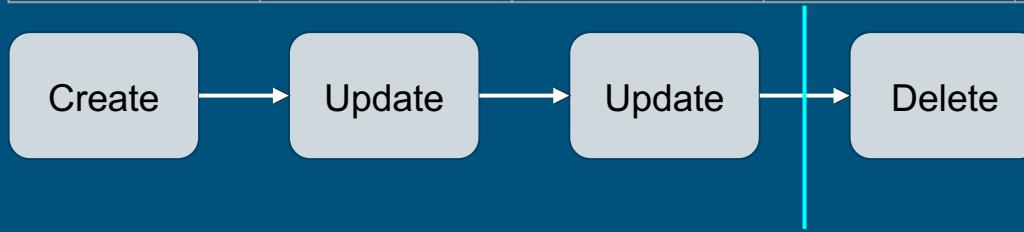
- Only one config file needed for a new Clickhouse table
- Init and keep syncing data in one app for a table
- Sync multiple data source to Clickhouse in minutes



BONUS

Time travel history state

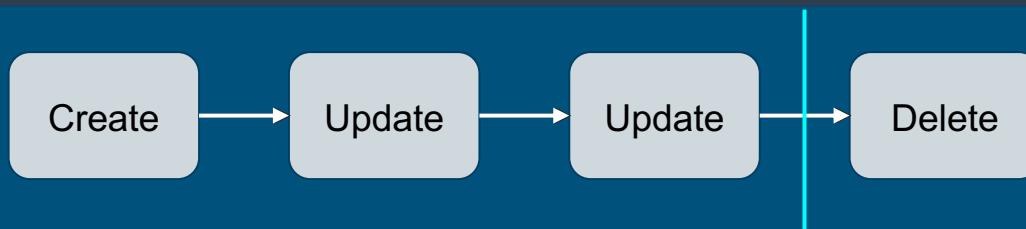
insert_id	id	name	flag	update_time
UUID0	2	Bob1	2	2019-10-01 00:00:00
UUID1	2	Bob1.5	2	2019-10-01 12:00:00
UUID2	2	Bob2	2	2019-10-02 00:00:00
UUID3	2	Bob2	3	2019-10-03 00:00:00



BONUS Time travel history state



```
WITH toDateTime('2019-10-02 12:00:00') AS target_time
SELECT argMin(insert_id , target_time-update_time) insert_id,
       argMin(flag, target_time-update_time) flag
FROM classroom_all
WHERE target_time-update_time>0
GROUP BY id
HAVING flag<3;
```



Future

- Auto configure through web
 - Auto deploy on Kubernetes
 - Open source?
 - Github: kevwan
-

Q&A



Thanks