

UDF in ClickHouse

Concept, Development, and Application in ML Systems

About CraiditX

CraiditX 氮信, a finance AI startup since 2015

The Main Business

- AI-based risk control
- AI-based marketing
- AI-based customer service
- ...

Our Partners and Customers



About Me

Chenzhang HU 胡宸章


R&D Engineer at CraiditX

Focusing on AI systems and algorithms

Active GitHub User

- <https://github.com/hczhcz>
- Interested in computer system and language stuff
- 8 organizations, 90+ repos, 600+ followers

ClickHouse Contributor



hcz
hczhcz

Coding for fun & +1s

[hczhcz.github.io](https://github.com/hczhcz)

Block or report user

Organizations



OLAP in ML Systems

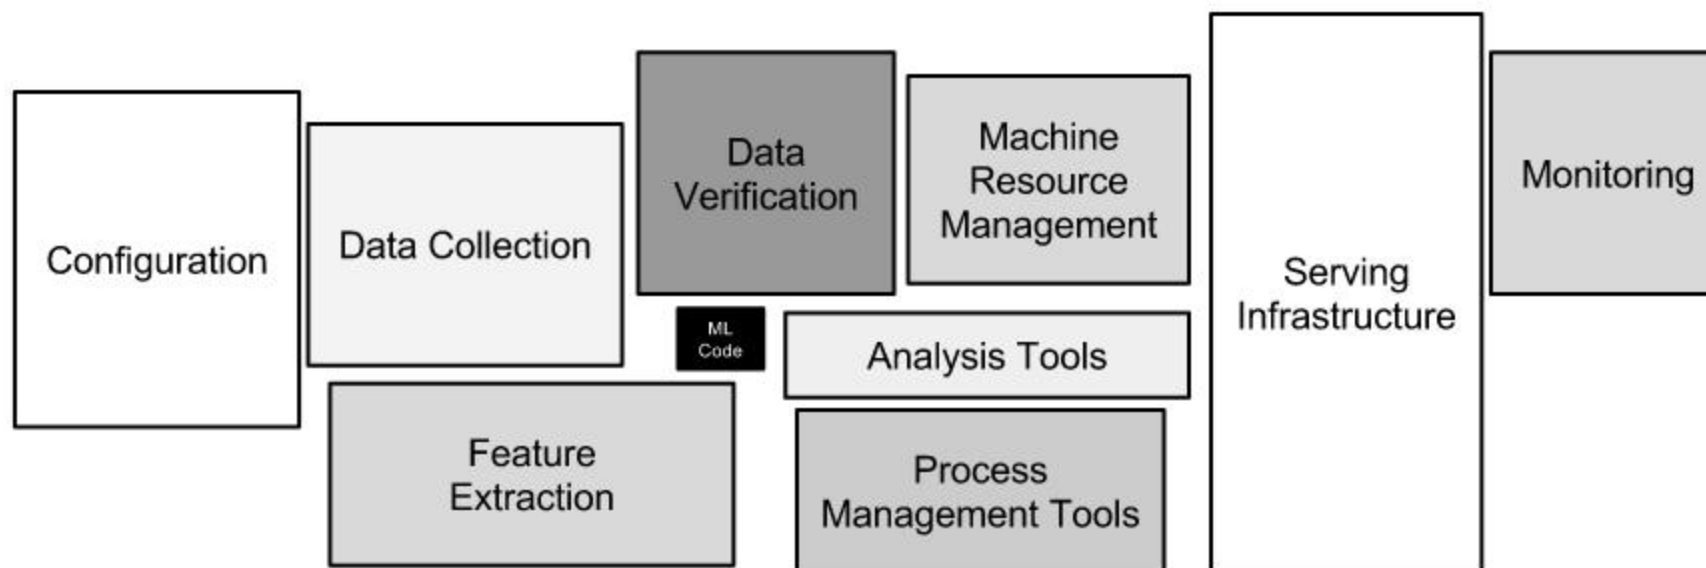


Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

Intensive Tasks in a ML System

- Pre-analyzing the data
- Extracting features
- Constructing relationship graphs
- Generating reports
- ...

Intensive Tasks in a ML System

- Pre-analyzing the data
 - = Finding the useful part of data + Summerizing data
- Extracting features
 - = Joining data + Grouping data + Doing math or pattern matching
- Constructing relationship graphs
 - = Identifying connections + Grouping data into source-destination pairs
- Generating reports
 - = Joining data + Summerizing data
- ...

The data processing scenario is very similar to OLAP

A Database is not Just a “Database”

What an English Dictionary Tells You

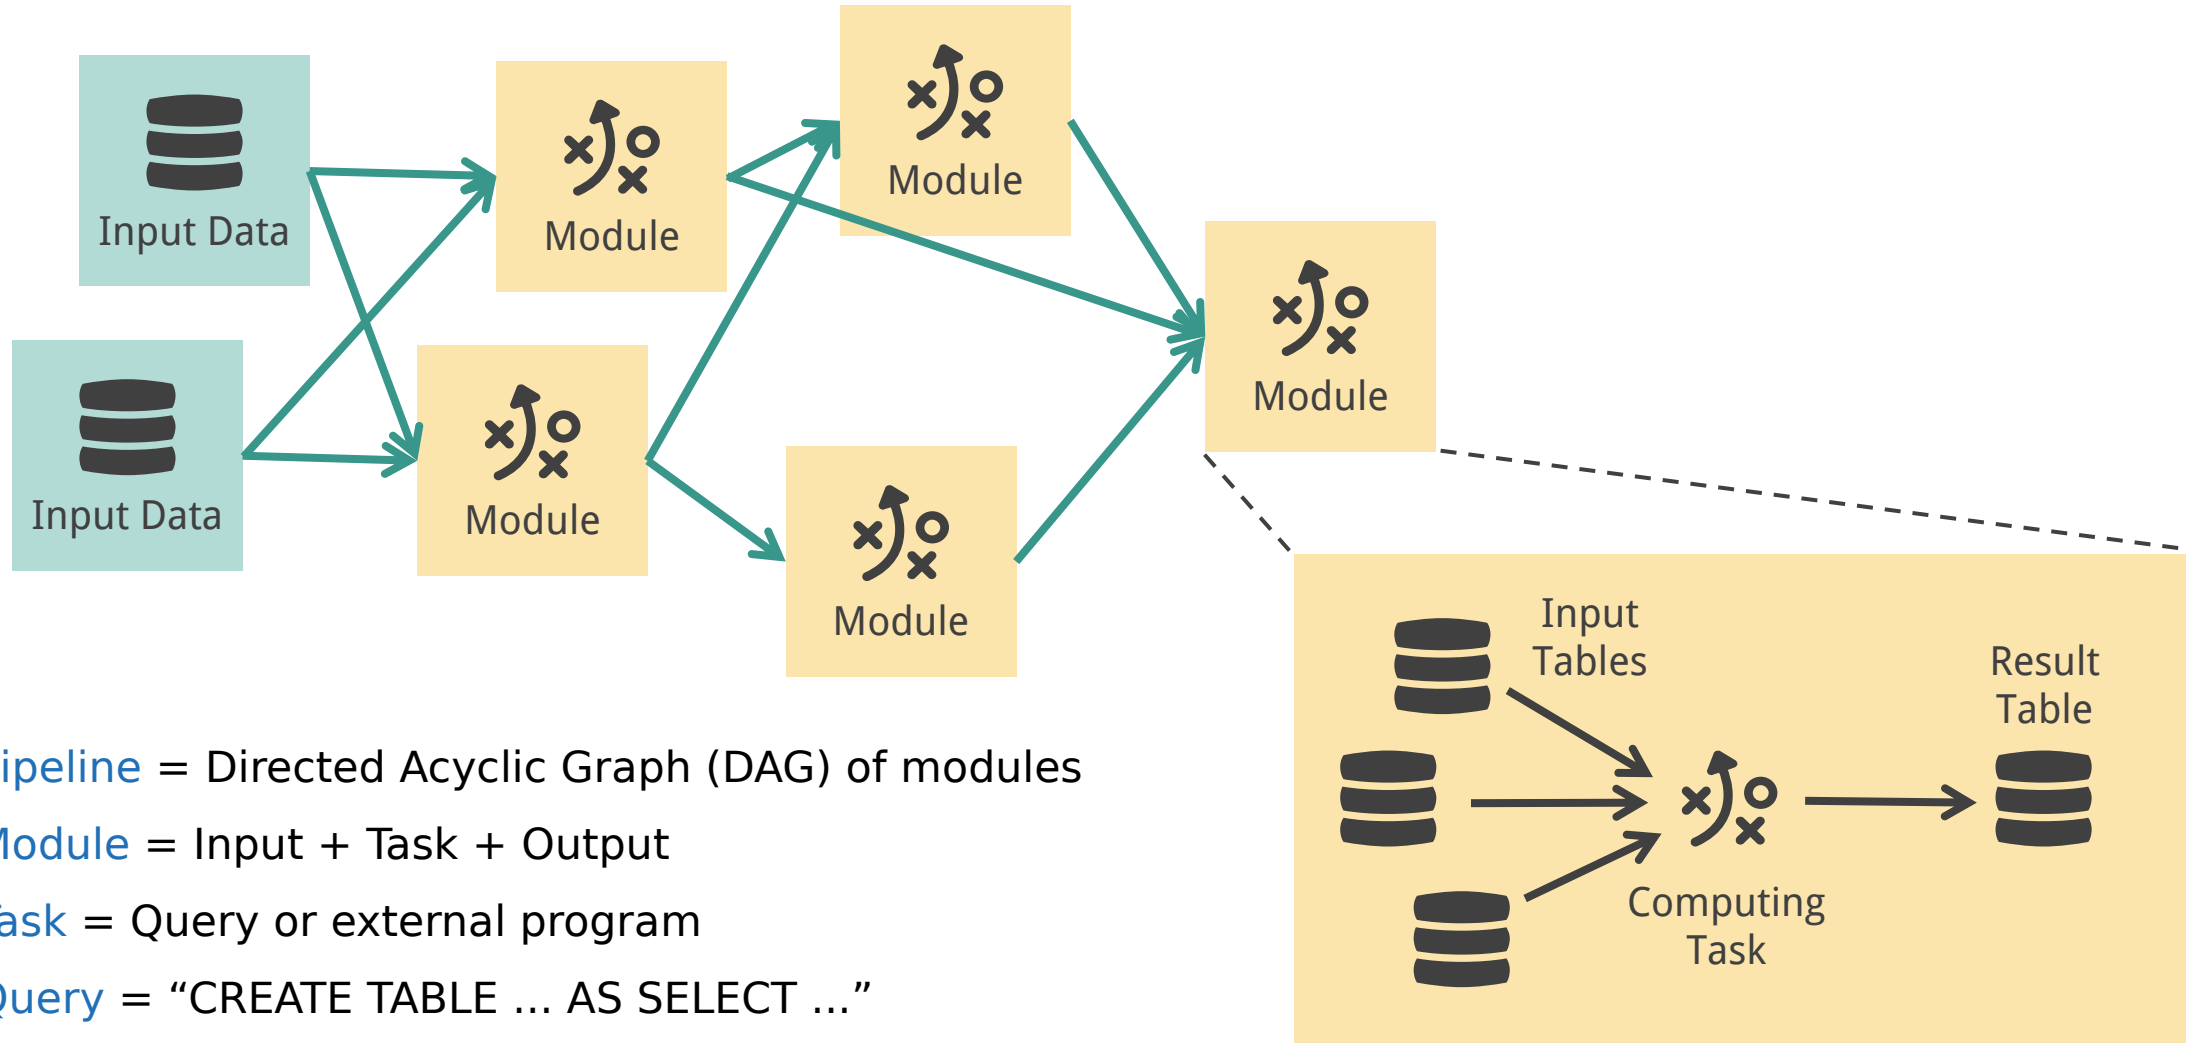
- database /'deɪtə,beɪs/

A collection of data **stored** in a computer that can easily be used and added to

What We Actually Do



A Database System and A ML Pipeline



Why ClickHouse

Limited hardware resources & time → efficiency matters

Performance

- Each node is able to handle [billions of rows](#)
- Most queries can be done in [0.1s-10min](#)

Ease of Use

- SQL
- Portable binary deployment with few dependencies

Customization

- The straight-forward code structure and the well-designed API
- We maintains a [custom build](#)



The UDF Magic

When the “Standard” SQL is not Enough

Functionality Limitations

- Scanning data back and forth
Example: Recognizing a behavioral pattern in the time series data
- Or even... randomly
Example: Finding a shortest path in the graph
- Iterating
Example: Training a regression model
- Handling domain-specific data
Example: Computing the edit distance between two strings
- ...

When the “Standard” SQL is not Enough

Performance Concerns

- A “dull” SQL solution can be 10x-1000x slower than a native C++ program

Example: Multiple self-joining on time series

Ease of Use and Maintainability

```
SELECT skewPop(x) FROM data
```

```
SELECT centralMoment(3)(x) / pow(stddevPop(x),  
3) FROM data
```

```
SELECT (sum(pow(x, 3)) / count() - 3 *  
sum(pow(x, 2)) * sum(x) / pow(count(), 2) + 2 *  
pow(sum(x), 3) / pow(count(), 3)) /  
pow(stddevPop(x), 3) FROM data
```



A UDF Can Do the Trick

User-Defined Functions (UDF)

- A function provided by the user

UDF in ClickHouse

- Scalar functions
- Aggregate functions & combinators
- Table functions & storage engines

Usage Examples in Our ML Systems

Data Preprocessing

Filling invalid date strings in a time series

Feature Engineering

Calculating average values within a window

Connection Recognition

Finding persons with similar street addresses

Zoo of Our UDF

Windowed Aggregate Functions

```
SELECT windowRefer(30)(date, value) FROM data
```

```
SELECT windowReferEx(30, 'quantile(0.2)')(date, value) FROM data
```

- Computing the aggregate results of the rows within each time window
- When a row enters, it will be “added” to the aggregation state
 - Then, the aggregate function yields a result field
- When a row leaves, it will be “removed” from the aggregation state
 - It requires some special algorithms and tricks
 - Most of the aggregate functions and combinators are re-implemented
- Typical usage: Feature computing

A further idea: -Window combinator

Funnel Automata Functions

```
SELECT yFunnel(6, '
    rule(r1)(require(get(1) = ''index.html'')),
    rule(r2, r1)(require(get(1) = ''product_info.html''),
        product_id(get(2)), return(''viewed'')),
    rule(r3, r2)(require(get(1) = ''product_buy.html''),
        verify(product_id = get(2)), return(''purchased''))
')(timestamp, page, product_id)
FROM pageview
```

- Matching behavior sequences within time window
- Inspired by Analysys OLAP Challenge 2018 (Funnel Analysis)
- Featuring built-in [automata description DSL](#)
 - It is implemented using the parser facilities in ClickHouse

Array Functions

SELECT arraySplit(x -> x >= 10, [11, 4, 5, 14]) = [[11, 4, 5], [14]]

SELECT arrayFill(x -> x > 0, [1, 2, 0, 0, 3, 0]) = [1, 2, 2, 2, 3, 3]

- Handling [time series data](#)
 - With arrays, the data block boundary will not affect the result
- Based on the common implementation of high-order array functions
 - FunctionArrayMapped.h
- Used in Analysys OLAP Challenge 2019 (Session Analysis)
- Pull request [#7294](#) and [#7380](#)

Integration of Third-party Libraries

FuzzyWuzzy

- Levenshtein distance (edit distance)
 - By character
 - By token
- Popular among ML algorithm engineers

Simdjson

- An [extremely fast](#) JSON parser based on AVX2 Instruction Set
- Structured data extraction (JSONExtract)
 - We can pass the type as a parameter just like in CAST function
- Difficulties in cross-platform compatibility
- Pull request [#4686](#) and [#5124](#)

Miscellaneous

Statistics

- Simple linear regression
- Skewness and kurtosis
- Weight-based entropy
- Pull request [#4668](#) and [#5200](#)

Aggregate Function Combinators

- -OrDefault and -OrNull
 - When an aggregate function has **nothing to aggregate**
- -Resample
- Pull request [#5590](#) and [#7331](#)

UDF Development Explained

General Steps of UDF Development

Design the Interface

- Meta information

Example: Will the return value change over time?

- Arguments and the return value

Create the Function Body

- Scalar functions: Handling data blocks
- Aggregate functions: Maintaining internal states, adding, merging, and reducing

Put Things Together

- Registering your function
- Build and run!

Type Checking and Inference

What should we consider? Let's see...

```
arraySum([1, 2, 3])
```

- Number of argument = 1
- Data type of the only argument = Array(X)
 - X should be UInt, Int, or Float
- Data type of the return value = Extended(X)
 - We prefer the largest native numerics, i.e. UInt64, Int64, or Float64
- Data types vs. column types
- Furthermore?
 - The real arraySum also supports lambda

Scalar Function Implementation

Then...

```
arraySum([1, 2, 3])
```


- Data are handled **per block**
- Column type of the argument = ColumnArray(ColumnVector(X), ColumnOffset)
 - ColumnConst can be handled explicitly or automatically

value	offset	arraySum	
[1, 2, 3]	3	6	<- sum of value[0..2]
[9, 10]	5	19	<- sum of value[3..4]
[4, 5, 6]	8	15	<- sum of value[5..7]

Aggregate Function Implementation

`avg(x)`

- Data are handled per group and row
- Internal state = (`count()`, `sum(x)`)
 - To add one row
 - `sum(x) += x[i]`
 - `count() += 1`
 - To calculate the result
 - `avg(x) = sum(x) / count()`
 - Also, we need to implement serialization, deserialization, and merging
- Some aggregate functions require more [advanced algorithms](#)
 - `uniq(x)`, `median(x)`, ...
 - Most of them depend on sampling tricks and/or special data structures



Going Further

Inline C++ in SQL

```
SELECT udsf('
    std::string udsf(std::string s)
    {
        return "hello, " + s;
    }
', 'world')
```

- Compiled and linked **dynamically**
 - It uses the embedded compiler facilities in ClickHouse
- Type inference driven by C++ Template Metaprogramming
- Also supported: Block-based scalar functions, aggregate functions

Okay for proof-of-concept usages, but ...

Zora: High-performance Algorithm Implementation Framework

Main Concepts

Column-oriented & Memory Densed

High memory efficiency and avoiding unnecessary IO
Smooth integration with ClickHouse, NumPy, Pandas, ...

```
template <typename pos_t, typename value_t,
void personalized_page_rank(
    pos_t &member_n,
    pos_t *&member_node,
    value_t *&member_rank,

    const pos_t *node_edge_begin,
    const pos_t *node_edge_end,
    const value_t *node_weight,
    const pos_t *edge_node_to,
```



Minimalism Fundamental Algorithm Components

Carefully chosen algorithm targeted at supporting the ML pipeline
Featured: Data structures, graph algorithms, statistical operators, ...

```
n, node, rank = zora.personalized_page_rank(
    df_node.edge_begin,
    df_node.edge_end,
    df_node.weight,

    df_edge.node_to,
    df_edge.weight,
```



Write Once, Available Everywhere

Native C++ Implementation with no third-party dependency
Auto-generated interface to Python3 and ClickHouse (UDF)

```
SELECT
    personalizedPageRank(node_edge_begin, node_weight) as rank
FROM (
    SELECT
        groupArray(edge_begin) as node_edge_begin,
        groupArray(edge_end) as node_edge_end,
        groupArray(weight) as node_weight
    FROM nodes
```



Questions / Comments